

CIBC Hackathon: Build Phase

Build a working version of your design: all four layers working end to end, with your chosen use case built in depth.

Deadline: **9:00 PM IST, 4 Oct.**

Timeline

When	What
3 Oct, 9:00 AM	Build phase starts, dataset released
3 to 4 Oct	Industry reps review the top 10 designs and run an insight session
4 Oct, 7:00 PM	Extra benchmark questions released (see Benchmark answers)
4 Oct, 9:00 PM	Submission deadline
After judging	10 finalists selected, get several days to refine
To be announced	Final presentations at the partner office

The data

Download: huggingface.co/datasets/nuxsh/maple-collections-hackathon

```
pip install -U huggingface_hub
hf download nuxsh/maple-collections-hackathon --repo-type dataset --local-dir
maple_data
cd maple_data && unzip maple_collections_release.zip
```

You need about 10 GB of free disk space (16 GB with the voice files). The `hf` command needs a recent `huggingface_hub`, so run the `pip install -U` line first.

File	Size	What
maple_collections_release.zip	2.3 GB	Start here. 31 tables (CSV and Parquet), call transcripts (JSON), policy documents, data dictionary, 500 public labels each for notes and transcripts
maple_collections_voice.zip	2.4 GB	Optional. 2,000 call recordings (WAV)

- Read `maple_collections_release/README.md` first. It covers file conventions, loading code and the customer key each source system uses.
- One fictional bank (Maple Bank), 1,000,000 customers, Oct 2016 to Sep 2026, snapshot date 28 Sep 2026. All data is synthetic.
- Data quality issues are deliberate: duplicates, missing values, mixed date formats, mismatched IDs, schema drift.
- Use DuckDB or Polars for the large tables. Loading everything into pandas will be slow.
- Protected attributes (gender, marital status, citizenship, household, newcomer, accessibility, vulnerability, accent) are included only to test fairness. **Do not use them in decisions.**

What to build

Every layer must work at a basic level, connected end to end. Go deep on your Layer 4 use case.

Layer	Must show
1. Data Product Factory	Raw to curated to golden C360 pipeline (code). Customer ID matching across the five source systems, with a match confidence. Data quality checks.
2. Insight and NLP	Plain-English questions answered over the C360. The SQL or source shown with every answer. Out-of-scope questions refused.

Layer	Must show
3. Feature Store	Features from structured data and from text or voice. One definition used for both training and live scoring.
4. Models and Decisioning	A working model or agent for your use case. A plain-English explanation of each decision. A point where a human reviews or overrides.

You may change your design from the system design phase. Say what changed and why in your README.

What to submit

Submit through the form by 9:00 PM IST, 4 Oct. Put everything except the video and deck in your repo.

1. **Code repository.** Private GitHub repo with the organisers added as collaborators (this will be clarified further into the hackathon). Commit throughout the hackathon. The README must cover: how to run it, what changed from your design, and which AI tools you used and for what.
2. **Architecture diagram.** Final version, as a PDF or PNG in the repo.
3. **Data contract.** YAML or JSON file for your golden C360: owner, schema, quality rules, allowed uses, refresh. Use the example contract DC-COLL-001 in files/docs/data_contracts/ as a starting point.
4. **Data quality report.** The issues you found in the data, how many records each affected, and how you handled them. Markdown or PDF.
5. **Benchmark answers.** A CSV with one row per question, in the format of the answers template in the release. The release README's "Benchmark questions" section and the dev answers file show how to fill it. Columns:
 - `question_id`
 - `answer` (your system's answer)
 - `sql_or_sources` (the SQL it ran, or the documents it used)
 - `refused` (true or false)

Include every question in the released benchmark set, plus the extra questions released at 7:00 PM on 4 Oct. Answers are scored automatically. Run your system as it is; don't edit answers by hand.

6. **Demo video.** Up to 5 minutes, unlisted YouTube or Google Drive link. Show your system working on real data from the dataset.

7. **Pitch deck.** Up to 10 slides, as a PDF.

Judging

Criterion	Weight
Technical quality and working demo	30%
Business understanding	20%
Presentation and storytelling	20%
Innovative ideas	15%
Team work	15%

Your benchmark answers score counts towards Technical quality and working demo. Judges will also ask unseen questions live during review.

Rules

- Commits made after 9:00 PM on 4 Oct are not considered.
- Every decision shown to a collections agent comes with a plain-English reason.
- A human can review anything that affects a customer.
- No protected attributes in decisions. Flag financial hardship.
- Paid LLM APIs are allowed. The data is synthetic, so you can send it to any API. API costs are your team's.
- Report problems with the dataset to the organisers on the doubts channel.