

CIBC Hackathon - Design Phase Rubric

How to score

Score each criterion from 1 to 5, then compute the weighted total out of 100. The top 50 go to the build phase.

Criterion	Weight
Business impact	25%
Technical quality	25%
Smart use of AI	20%
Clarity	20%
Governance and responsible AI	10%

Gates

Check these first.

Problem	Action
Submitted after 11:00 PM	Not scored
Clear AI Slop	Up to scorer
Not following format	Up to scorer
No architecture diagram	Clarity and Technical quality capped at 2
Over 6 pages	Score only the first 6
No Layer 4 use case chosen	Business impact capped at 2
No data contract	Governance capped at 2
Copied from another team	Both disqualified, pending review

Judge rules.

- Score from the PDF only. Don't infer/assume/fill in the gaps.
- Fields marked "to be confirmed against the data" are fine.
- Score each criterion on its own.

Business impact (25%)

Does the team understand collections, and would their use case actually help?

Look for:

- A named user (collections agent, team lead, strategy analyst), not "the bank"
- A specific problem that user has today
- One metric the solution moves: promise-kept rate, roll rate, cure rate, cost per contact, right-party contact rate
- How they would measure the change (A/B test, before and after, holdout group)
- Use case fits the problem: e.g. routing for a capacity problem, next best action for a treatment problem

Score	What it looks like
1	Generic "use AI to improve collections". No user, no metric.
3	Clear use case and user. A metric is named but not tied to how it would be measured, or the link from feature to metric is vague.
5	Specific user and pain point, one clear metric, a credible way to measure it. Shows real collections understanding (buckets, promises to pay, hardship, contact rules).

Technical quality (25%)

Is the architecture complete, sound and buildable? Check each layer.

Layer	Strong design shows
1. Data Product Factory	Raw to curated to golden C360 stages. A concrete customer ID matching approach (rules on name, DOB, address, phone, plus a confidence score). How messy data is handled: duplicates, missing values, bad dates. Data quality checks.
2. Insight and NLP	NL-to-SQL over the curated or golden data, using the metric definitions. RAG over notes and transcripts. The SQL or source shown with every answer. A way to refuse out-of-scope questions.
3. Feature Store	Named features from structured data (e.g. days past due, broken promises) and from text or voice (e.g. hardship signal). One definition for training and live scoring. Refresh frequency.
4. Models and Decisioning	A model or agent that fits the use case. Inputs from the feature store. Output an agent can act on. How it is evaluated.

Also check: batch vs real-time split is sensible, named tools rather than generic boxes, and the 36-hour build plan is realistic.

Score	What it looks like
1	One or more layers missing. Boxes with no real components. Nothing on ID matching or messy data.
3	All four layers present, but one or two are thin or hand-wavy. Tools named. ID matching mentioned but not explained.
5	All four layers designed with clear components and data flow. ID matching and data quality handled concretely. Feature store reused across training and scoring. Build plan is achievable in 36 hours.

Smart use of AI (20%)

Is AI used where it adds value, with a reason each time? More AI is not better.

Look for:

- AI in the places the brief suggests: mapping and SQL generation (Layer 1), NL-to-SQL and RAG (Layer 2), features from text or voice (Layer 3), the decision model or agent (Layer 4)
- A reason for each use, and places where they chose not to use AI (e.g. simple rules for contact limits)
- How AI output is checked: LLM features tested against labels, NL-to-SQL checked against benchmark questions, guardrails against made-up answers
- Explanations of decisions an agent can read, not just a score

Score	What it looks like
1	"We use an LLM" with no detail, or AI forced into every box with no reason. No checking of AI output.
3	AI used in sensible places with some reasoning. Little on how outputs are validated or explained.
5	Each AI use is justified and the right tool for the job. Clear validation of AI output. Decisions come with plain-English reasons. Knows when rules beat models.

Clarity (20%)

Can a judge understand the design in 10 minutes?

Look for:

- Uses the required headings in order
- Summary states problem, use case and user in 3 lines
- Diagram is readable and matches the written design
- Specific to this problem: mentions the actual tables, buckets, promises to pay
- No padding or generic filler

Score	What it looks like
1	Hard to follow. Headings missing. Diagram unreadable or contradicts the text. Mostly generic text.
3	Follows the structure. Understandable with effort. Some padding or vague sections.
5	Easy to follow on first read. Diagram and text match. Every section says something specific.

Governance and responsible AI (10%)

Does the design follow the four ground rules in a concrete way?

Look for:

- **Data contract:** owner, schema, at least 3 quality rules, allowed uses, refresh
- **Explain every decision:** the agent sees why, not just what
- **Human in the loop:** a named point where a person reviews or overrides before anything affects a customer
- **Fairness:** protected attributes excluded, and some thought about proxies (e.g. postal area)
- **Hardship and vulnerability:** detected and routed to a human or hardship program
- **Access control:** PII masked or restricted by role

Score	What it looks like
1	Governance missing or one generic line. No data contract.
3	Data contract present. Human review and fairness mentioned but not placed in the flow.
5	All four ground rules handled concretely and visible in the diagram. Hardship handling and proxy risk addressed.

Red flags and selection

Red flags. Note them in the comments column. Reflect them in the relevant score.

- Generic text that could fit any hackathon (likely AI-written with no thinking behind it)
- Uses protected attributes, or features like age or postal area without a fairness check
- Fully automated actions on customers with no human review
- Claims that can't work: "100% accurate", real-time everything with no cost or latency thought
- Build plan that clearly won't fit in 36 hours
- Clearly AI Generated diagrams/graphics/slides

Tie-breaks, in order:

1. Higher Technical quality score
2. More realistic 36-hour build plan
3. Higher Business impact score